

ANN

Artificial Neural Networks



STUDENT 3

Test: 7/10

Grades: 6/10



QUIZ

Does the student get Accepted?

☐ Yes

☐ No

Acceptance at a University




QUIZ

Does the student get Accepted?

☒ Yes

☐ No

Acceptance at a University



BOUNDARY:

A LINE

$$w_1x_1 + w_2x_2 + b = 0$$

$$Wx + b = 0$$

$$W = (w_1, w_2)$$

$$x = (x_1, x_2)$$

$y = \text{label: 0 or 1}$

PREDICTION:

$$\hat{y} = \begin{cases} 1 & \text{if } Wx + b \geq 0 \\ 0 & \text{if } Wx + b < 0 \end{cases}$$

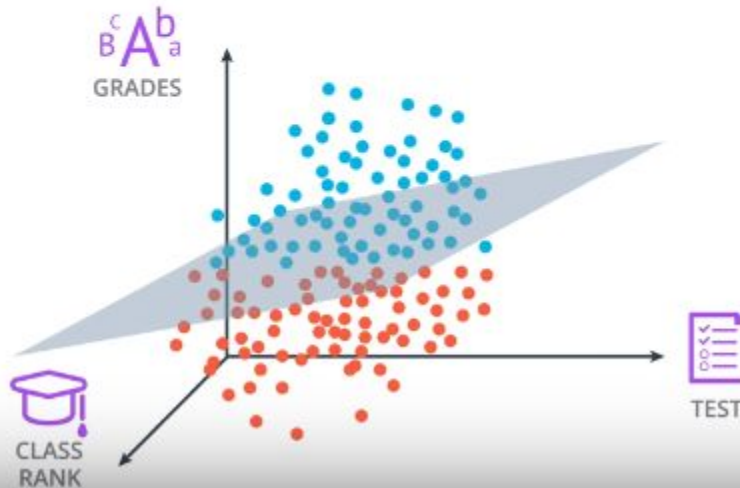
Applications of Deep Learning?



- Beating professional players at games like chess, checkers and go
- Detecting spam emails
- Predicting stock prices
- Classifying images
- Diagnosing illnesses
- Self-driving cars

What if we have more than 2 features?

Acceptance at
a University



BOUNDARY:

A PLANE

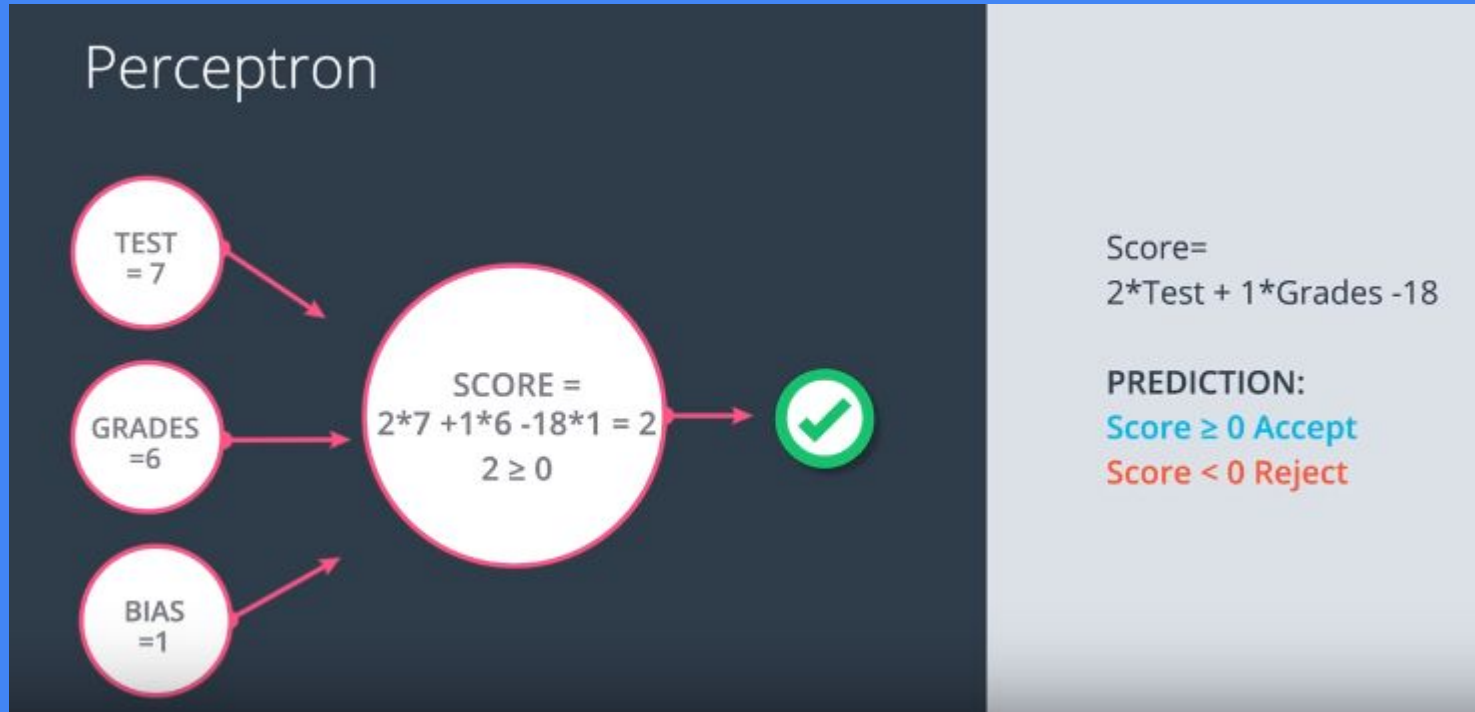
$$w_1x_1 + w_2x_2 + w_3x_3 + b = 0$$

$$Wx + b = 0$$

PREDICTION:

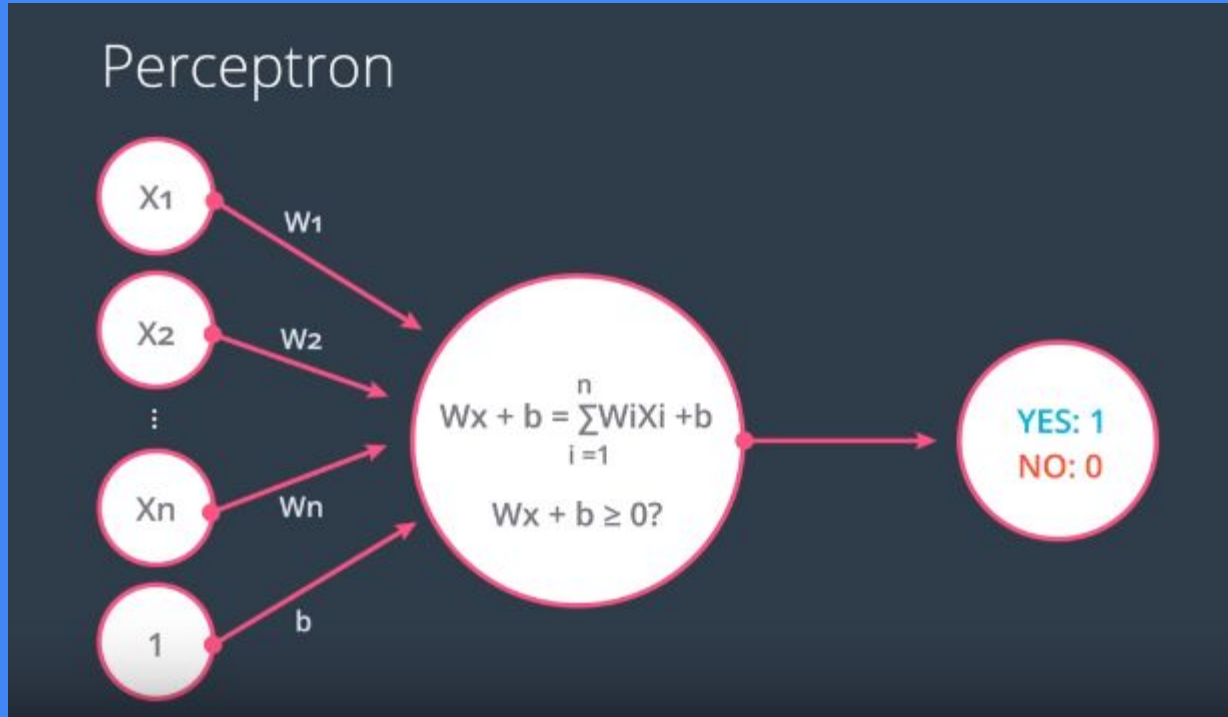
$$\hat{y} = \begin{cases} 1 & \text{if } Wx + b \geq 0 \\ 0 & \text{if } Wx + b < 0 \end{cases}$$

Let's suppose that the line demarcating the 2 classes is
 $2 \cdot \text{test} + 1 \cdot \text{grades} - 18$




The weights for the edges are 2,1,-18 respectively.

In general, this is how a perceptron model looks like!



We can separate the step of checking $Wx+b>0$





QUIZ QUESTION

Given Score = $2 * \text{Test} + 1 * \text{Grade} - 18$, suppose w_1 is 1.5 instead of 2. Would the student who got 7 on the test and 6 on the grades be accepted or rejected?

ANSWER : REJECTED!


Perceptron



Perceptron as logical operators

FOR AND OPERATION

X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1



What should be the weights for the perceptron?

ANSWER : weight1 = 1.0, weight2 = 1.0, bias = -1.5


FOR OR PERCEPTRON

OR Perceptron




weight1 = 1.0, weight2 = 1.0, bias = -0.5


XOR perceptron



$$Y = (AB)'(A+B)$$

XOR Multi-Layer Perceptron





LINE: $3x_1 + 4x_2 - 10 = 0$

POINT: (4,5)

LEARNING RATE: 0.1


3	4	-10
0.4	0.5	0.1
2.6	3.5	-10.1

NEW LINE

$$2.6x_1 + 3.5x_2 - 10.1 = 0$$

Learning rate


Big learning rate



It determines how fast or slow we will move towards the optimal weights.


- If it is **too big**, it maybe will not reach the local minimum because it just bounces back and forth.
- If it is **very small**, gradient descent will eventually reach the local minimum in a long time.

Small learning rate



Keep learning rate :neither too low nor too high.

Perceptron Algorithm



1. Start with random weights: w_1, \dots, w_n, b

2. For every misclassified point (x_1, \dots, x_n) :

2.1. If **prediction = 0**:

- For $i = 1 \dots n$
- Change $w_i + \alpha x_i$
- Change b to $b + \alpha$

2.2. If **prediction = 1**:

- For $i = 1 \dots n$
- Change $w_i - \alpha x_i$
- Change b to $b - \alpha$

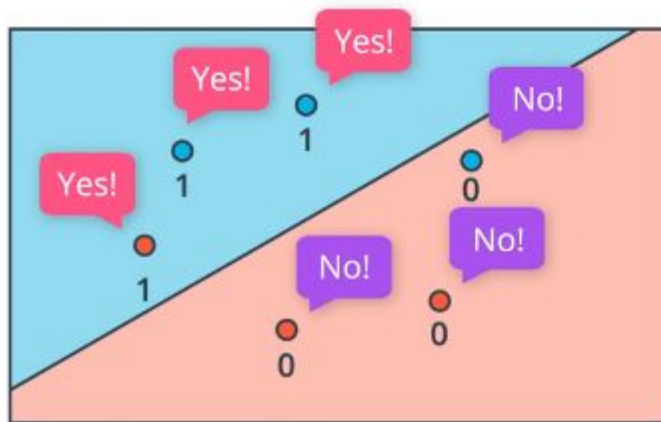
Error Function

The deviation of the actual output from the predicted output

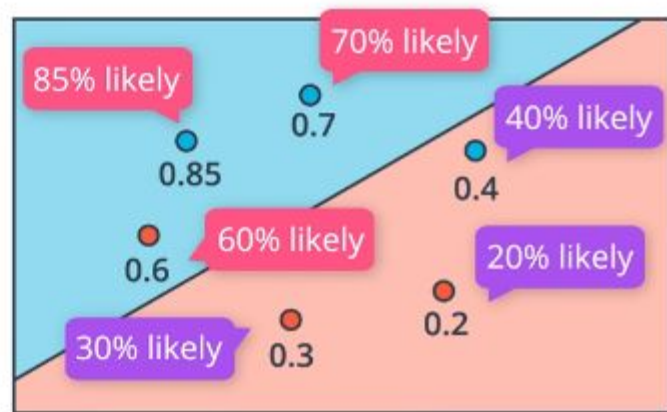
Tells the model how badly it is performing



Predictions




DISCRETE




CONTINUOUS

Activation Functions



DISCRETE:
Step Function

$$y = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$



CONTINUOUS:
Sigmoid Function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

The sigmoid function is defined as $\text{sigmoid}(x) = 1/(1+e^{-x})$. If the score is defined by $4x_1 + 5x_2 - 9 = \text{score}$, then which of the following points has exactly a 50% probability of being blue or red?

- a) (1,1)
- b) (2,4)
- c) (5,-5)

ANSWER: (1,1)

Multi class classification

Softmax Function


LINEAR FUNCTION

SCORES:


z_1, \dots, z_n

$$P(\text{class } i) = \frac{e^{z_i}}{e^{z_1} + \dots + e^{z_n}}$$

Probability



$$0.6 * 0.2 * 0.1 * 0.7 = 0.0084$$



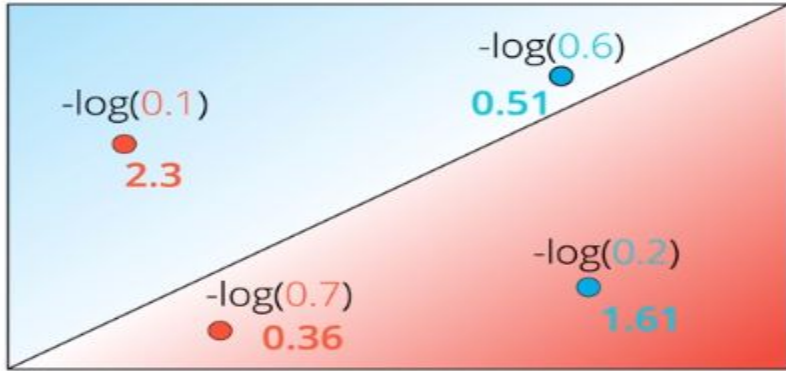
$$0.7 * 0.9 * 0.8 * 0.6 = 0.3024$$

What function turns products into sums?

- a) Sin
- b) Cos
- c) Log
- d) exp


ANSWER : LOG!

Cross Entropy



$$0.6 * 0.2 * 0.1 * 0.7 = 0.0084$$

$$-\log(0.6) - \log(0.2) - \log(0.1) - \log(0.7) = 4.8$$



$$0.7 * 0.9 * 0.8 * 0.6 = 0.3024$$

$$-\log(0.7) - \log(0.9) - \log(0.8) - \log(0.6) = 1.2$$

$$\text{Cross-Entropy} = - \sum_{i=1}^m y_i \ln(p_i) + (1 - y_i) \ln(1 - p_i)$$

🗑️ **X**

Logistic Regression

It is one of the most popular and useful algorithms in Machine Learning, and the building block of all that constitutes Deep Learning. It basically goes like this:

- Take your data
- Pick a random model
- Calculate the error
- Minimize the error, and obtain a better model

Gradient Descent

1. Start with random weights:

$$w_1, \dots, w_n, b$$

2. For every point (x_1, \dots, x_n) :

2.1. For $i = 1 \dots n$


2.1.1. Update $w'_i \leftarrow w_i - \alpha (\hat{y} - y) x_i$


2.1.2. Update $b' \leftarrow b - \alpha (\hat{y} - y)$

3. Repeat until error is small



What if the data is not linearly separable?

Neural Network






Neural Network



Feedforward Phase

Multi-layer Perceptron




PREDICTION

$$\hat{y} = \sigma \circ W^{(3)} \circ \sigma \circ W^{(2)} \circ \sigma \circ W^{(1)}(x)$$


ERROR FUNCTION

$$E(W) = -\frac{1}{m} \sum_{i=1}^m y_i \ln(\hat{y}_i) + (1 - y_i) \ln(1 - \hat{y}_i)$$



Backpropagation

Backpropagation



Multi-layer Perceptron



PREDICTION

$$\hat{y} = \sigma W^{(3)} \circ \sigma W^{(2)} \circ \sigma \circ W^{(1)}(x)$$

ERROR FUNCTION

$$E(W) = -\frac{1}{m} \sum_{i=1}^m y_i \ln(\hat{y}_i) + (1 - y_i) \ln(1 - \hat{y}_i)$$

GRADIENT OF THE ERROR FUNCTION

$$\nabla E = (\dots, \frac{\partial E}{\partial w_j^{(i)}}, \dots)$$

$$W_{ij}^{r(k)} \leftarrow W_{ij}^{(k)} - \alpha \frac{\partial E}{\partial W_{ij}^{(k)}}$$

Overfitting vs Underfitting

TYPE OF ERRORS



UNDERFITTING



OVERFITTING

Overfitting

CLASSIFICATION

ANYTHING BUT
DOGS THAT ARE
YELLOW, ORANGE
OR GRAY



DOGS THAT ARE
YELLOW, ORANGE
OR GRAY



TOO SPECIFIC

Underfitting

CLASSIFICATION


NOT ANIMALS





ANIMALS





TOO SIMPLE
UNDERFITTING



UNDERFITTING
(high bias)



JUST RIGHT




OVERFITTING
(high variance)

UNDERFITTING




JUST RIGHT



OVERFITTING



OVERFITTING



EPOCH 1

Training Error: BIG
Testing Error: BIG

EPOCH 20

Training Error: SMALL
Testing Error: SMALL

EPOCH 100

Training Error: TINY
Testing Error: MEDIUM

EPOCH 600

Training Error: TINY
Testing Error: LARGE

Goal: Split Two Points




QUIZ: WHICH GIVES A SMALLER ERROR?

- ☐ SOLUTION 1: $x_1 + x_2$
- ☐ SOLUTION 2: $10x_1 + 10x_2$

Goal: Split Two Points

QUIZ: WHICH GIVES A SMALLER ERROR?



Prediction: $\hat{y} = \sigma(w_1x_1 + w_2x_2 + b)$

○ SOLUTION 1: $x_1 + x_2$

Predictions:

$$\sigma(1 + 1) = 0.88$$

$$\sigma(-1 - 1) = 0.12$$

✓ SOLUTION 2: $10x_1 + 10x_2$

Predictions:

$$\sigma(10 + 10) = 0.9999999979$$

$$\sigma(-10 - 10) = 0.0000000021$$

Solution: Regularization

LARGE COEFFICIENTS \longrightarrow OVERFITTING

PENALIZE LARGE WEIGHTS

(w_1, \dots, w_n)


L1 ERROR FUNCTION $= -\frac{1}{m} \sum_{i=1}^m (1 - y_i) \ln(1 - \hat{y}_i) + y_i \ln(\hat{y}_i) + \lambda(|w_1| + \dots + |w_n|)$

L2 ERROR FUNCTION $= -\frac{1}{m} \sum_{i=1}^m (1 - y_i) \ln(1 - \hat{y}_i) + y_i \ln(\hat{y}_i) + \lambda(w_1^2 + \dots + w_n^2)$


SPORTS



DROPOUT



DROPOUT




EPOCH 1

EPOCH 2


EPOCH 3

EPOCH 4

GRADIENT DESCENT




RANDOM RESTART




Vanishing Gradient

SIGMOID FUNCTION




Solution :

HYPERBOLIC TANGENT FUNCTION



$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

RECTIFIED LINEAR UNIT (ReLU)



$$\text{relu}(x) = \begin{cases} x & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$

Batch Gradient Descent

Batch Gradient Descent **calculates the error for each example** within the training dataset and after all training examples have been evaluated, the model gets updated. This whole process is like a cycle and called a training epoch.

Advantages

- Computational efficient
- It produces a stable error gradient
- A stable convergence

Disadvantage

- Requires a lot of training epochs
- It also requires that the entire training dataset is in memory

Stochastic Gradient Descent

Stochastic gradient descent **calculates error for each training example** within the dataset. It updates the parameters for each training example, one by one.

Advantage

- The frequent updates results in a detailed rate of improvement.
- Requires less number of training epochs.

Disadvantages

- Computationally expensive
- The frequency of those updates can also result in noisy gradients, which may cause the error rate to jump around, instead of slowly decreasing.
- It might not reach global minimum

Mini Batch Gradient Descent

It is a **combination of Stochastic Gradient Descent and Batch Gradient Descent**. It splits the training dataset into small batches and performs an update for each of these batches.

- It creates a balance between the robustness of Stochastic Gradient Descent and the efficiency of Batch Gradient Descent.
- Common mini-batch sizes range between 50 and 256.
- It is the most common type of gradient descent algorithm.